

Dynamic Context Runtime: Bounded Attention over Unbounded History

Separating unbounded history from bounded attention with a representation ladder, a provenance graph, and a budgeted context planner

Stephan Botes · Cyber Sec · cybersec.org.za

Specification, implementation and experiments: github.com/s-b-repo/subnext

ABSTRACT

Language models degrade as irrelevant, stale, and superseded material accumulates in their context — a failure commonly called *context rot*. Enlarging the context window does not remove it: attention cost grows quadratically, mid-context material is used unreliably, and a corrected fact does not stop competing with the value it replaced. We describe the **Dynamic Context Runtime** (DCR), a system that separates *unbounded history* from *bounded attention*. History is stored as immutable, addressable spans and compiled into a typed dependency graph of facts, decisions, and derivations, each bound to the spans that ground it. Each turn, a planner assembles a small working set by solving a multiple-choice knapsack over (node, representation-level) pairs under an explicit token budget, so the model never sees the transcript — only the cheapest sufficient representation of what the current question needs.

On a lexically varied corpus the working set holds at 4.19 million tokens of history, 48,651 state nodes and 7/7 probes, with 235 tokens per query. We give a zero-dependency Rust reference implementation (13,721 lines in `src/`) that enforces the design's invariants in code rather than in prose: unsourced facts are rejected at insertion, nothing is deleted, and no stale node reaches the model. We evaluate it with a deliberately non-neural instrument — a deterministic line-matching reasoner shared by all conditions — which isolates the quality of *context assembly* from the quality of the model. On a 300-turn synthetic incident transcript (27,362 tokens), DCR answers all seven probes using 462 tokens per query on average, 59× less attention than the full history, and answers all seven from a 200-token budget. Across a 33× growth in history (8.5k → 283k tokens) the working set stays flat at ~410 tokens with no loss of correctness. An ablation identifies which mechanisms carry which probes — and reports two that carry nothing on this corpus. A positive control turns the runtime's stale-fact metric from a zero that could not fire into a measured guard, and a read-coverage metric — the offline dual of audit completeness — shows the fraction of history the model ever actually sees collapsing toward zero as history grows: the blind region, not the working set, is what grows with N . Because durable memory is only auditable if it can prove it returned what it stored, context is held in a content-addressed, append-only container under a Merkle root, with chained generations and anti-rollback; an adversarial probe confirms that object corruption, checkpoint rewriting, and rollback are each detected, and we state precisely where the guarantee stops — tamper-evident, not tamper-proof. Against a second corpus built so that similarity is misleading and refusing is sometimes correct, **DCR scores 2/5 and loses to recursive map-reduce**; we report the three failures, since a benchmark incapable of showing the system losing is not evidence. We state plainly what the evaluation cannot show: it is a synthetic corpus, the instrument is not a language model, and we do not compare against a recursive-language-model implementation.

1 Introduction

A long conversation is not a long document. It contains restatements, digressions, values that were later corrected, and reasoning whose conclusions have long since been settled. A model asked a question at turn 300 must find what matters inside all of it. Practice has converged on two responses: make the window bigger, or retrieve into it. Neither addresses the underlying problem, which is that *the transcript is being used as the working memory*.

Bigger windows carry a quadratic attention cost in sequence length N [1], and empirically models use material in the middle of a long context less reliably than material at

either end [2] — so the marginal token added to a window is not merely expensive, it can be actively harmful to what is already there. Retrieval-augmented generation [3] reduces the volume but collapses the problem to a single representation and a single signal: chunks, ranked by similarity. It has no notion that a retrieved chunk was superseded three turns ago, and no way to hand back an exact byte range when the question is "quote it verbatim" rather than "what is the gist".

A more recent framing treats the context as a program variable: the model inspects, slices, and recursively queries it in a REPL rather than reading it as a flat prompt.¹ This is the right primitive — context as a manipulable object rather than a wall of text — but it relocates the problem rather

than removing it. Every recursive subcall still reconstructs what it needs from raw text: re-reading the same spans, re-deriving the same conclusions, and stalling on slicing and compression at each level of recursion.

Our position is that the missing layer is a *runtime*: a system that owns unbounded external state and hands the model a small, deliberately planned working set. The slogan is **unlimited history $\neq$ unlimited attention**, and the design consequence is that context assembly should be an explicit, budgeted optimisation problem rather than an emergent property of a prompt template.

Contributions.

1. A system design (§3) in which the same material exists simultaneously at four cost levels, facts are typed graph nodes bound to immutable source spans, revision is recorded rather than applied destructively, and the working set is chosen by an exact multiple-choice knapsack under a declared token budget.
2. A zero-dependency Rust reference implementation (§4) in which the design's stated invariants are enforced by the type system and by runtime checks, not merely asserted in documentation.
3. An evaluation methodology (§5) that isolates context assembly from model quality by holding the reasoner fixed and non-neural across all conditions, together with four measurements: a context-rot comparison, a scaling result showing the working set flat under 33 $\times$ history growth, a budget sweep, and an ablation that names the load-bearing mechanisms — including two that are not load-bearing on our corpus.
4. An explicit account (§6) of what this evaluation cannot establish.

2 Background and related work

2.1 Long context and its failure modes

Self-attention is quadratic in sequence length [1], motivating a long line of efficiency work — sparse and sliding-window attention [4], IO-aware exact attention [5] — that reduces the constant or the asymptotics of *computing* attention. That work is orthogonal to ours: it makes a large window cheaper to process, while we argue about what should be in the window at all. The positional reliability problem [2] is not addressed by making attention cheaper.

2.2 Retrieval-augmented generation

RAG [3] retrieves passages by learned or lexical similarity [6,7] and conditions generation on them. It is the closest deployed analogue to what we propose and differs on three axes. First, *unit*: RAG returns document chunks; DCR returns whichever representation is cheapest and sufficient, which may be an exact span, a structured fact, or a computed result. Second, *signal*: RAG ranks by similarity;

DCR additionally traverses typed dependency edges, so justifying a decision is a graph walk rather than a hope that the evidence happens to be lexically similar to the question. Third, *state*: RAG's corpus is static with respect to the conversation, whereas DCR's store is written by the conversation and must therefore handle correction, contradiction, and staleness as first-class events.

2.3 Agent memory systems

Generative-agent architectures maintain a memory stream scored by recency, importance, and relevance [8], and MemGPT [9] frames context management as virtual memory with paging between a fixed window and external storage. DCR shares the paging intuition and differs in what is paged and how the decision is made: the unit is a typed, provenance-bound node rather than a text fragment; the same node may be paged in at four different costs; and the selection is the solution to a stated optimisation problem rather than a heuristic score threshold. We also make invalidation explicit — a memory whose supporting fact changed is marked stale and excluded, rather than remaining eligible for retrieval at a lower score.

2.4 What databases already settled

Most of what a memory runtime does is not novel; it is a database problem that the language-model literature has been re-deriving. Provenance — the ability to say *where* a value came from and *why* it holds — is a well-studied property with formal characterisations [10]. Deciding when a cached derived value must be recomputed after its inputs change is materialised-view maintenance [11]. Reading a consistent snapshot while a background writer mutates the store is snapshot isolation [12]. Choosing a maximum-value subset under a capacity constraint, where each item has several mutually exclusive variants, is the multiple-choice knapsack problem [13]. Speculatively materialising state before it is requested, and discarding the work when the prediction was wrong, is prefetching [14]. DCR is in large part an argument that these should be imported wholesale rather than approximated.

3 Design

3.1 Two systems, asymmetric

DCR splits the agent into a **Reasoner** and a **Memory Runtime**. The Reasoner is a language model with a small, high-attention working set; it owns the current computational state and none of the history. The Memory Runtime owns everything ever seen and decides which representation of it to return. The asymmetry is the point: almost everything the Memory Runtime does — span addressing, index lookup, dependency traversal, memoisation, invalidation — requires no model inference at all. In our implementation the Memory Runtime makes zero model calls.

MEMORY RUNTIME




Figure 1. Per-turn data flow. Ingestion addresses raw bytes into immutable spans and extracts typed state; higher representation levels are built lazily. A query is classified, seeded from the index, expanded along dependency edges, assigned levels, and bounded by the attention budget. The model sees only the resulting working set, and may escalate a node to raw bytes if the compact form proves insufficient.

3.2 The representation ladder

The same material is available at four levels of cost and fidelity:

LEVEL	PAYLOAD	TYPICAL USE
L0	exact bytes of a span	quote, identifier, legal or audit fidelity
L1	extractive summary	gist of a large span
L2	structured state object	value lookup; the fact cache
L3	computed result	a value that is a function of known inputs

Only L0 addressing is eager. L1, L2, and L3 are built on first demand and cached on the node, so ingesting a large document is cheap until something asks about it.

L2 deserves a clarification, because "semantic vector" is not something a model can read. In DCR the vector is L2's *index key* and a structured state object is its *payload*. Admitting a node at L2 means the model sees

```
server.ip = 10.0.9.7 · conf=0.90 ·
spans=s_9d6230ba
```

rather than the paragraph that sentence came from. This is the level at which a long transcript collapses into a few dozen objects, and it is why compression here is safe: the object carries pointers back to the exact bytes, so any loss is recoverable by escalation.

3.3 A typed graph with mandatory provenance

State is a directed graph of typed nodes — **evidence**, **claim**, **calculation**, **decision**, **goal**, **constraint**, **open_question** — connected by typed edges: **supports**, **derived-from**, **contradicts**, **supersedes**. Every node

carries source spans, dependencies, a confidence, and a timestamp.

Three invariants are enforced at insertion rather than documented:

- Nothing enters ungrounded.** Every non-evidence node must reach at least one evidence node through its dependencies. A cached fact with no spans is a hallucination with a database row, and the insertion path rejects it.
- Nothing is deleted.** Revision is a **supersedes** edge and a status change; conflict is a **contradicts** edge with both sides retained. The transcript stays a valid audit log because no correction ever rewrites history.
- No stale node reaches the model.** When a node's content changes, its transitive dependents are marked stale and excluded from planning until revalidated.

Invalidation cascades are *bounded*. A single small correction can otherwise mark an arbitrarily large subgraph stale and stall the turn, so the runtime marks eagerly up to a cap and defers the tail to a lazy queue drained before the next plan.

3.4 Correction as a first-class event

The mechanism that most directly attacks context rot is what happens when a new claim disagrees with a live claim on the same subject key. The runtime records a **contradicts** edge, then — by default — **supersedes** the older node, which marks its dependents stale. The superseded node remains in the graph and is never planned into the window again.

Two refinements proved necessary in practice. First, evidence spans that sourced a superseded value are annotated rather than hidden (NOTE=contains a value corrected later; current value in ...): deleting them

would break the audit log, but admitting them unmarked is exactly how a settled contradiction walks back into the window. Second, ingest order is only approximately chronological — files read from a directory, transcripts merged from two sources — so an explicitly corrective statement is protected: ordinary material arriving after it records a contradiction but does not supersede it, and both sides enter the window flagged for adjudication.

3.5 The attention budget

Given a budget B in tokens, the runtime solves

$$\begin{array}{llll} \text{maximise} & U(S) & \text{over } S \subseteq M & \text{subject to} \\ & \sum_{x \in S} \text{cost}(x) \leq B \end{array}$$

where M is the candidate set. Because the same node may be admitted at several levels, and at most one, this is a multiple-choice knapsack [13]. We solve it exactly by dynamic programming over quantised costs, with the quantum chosen so the table stays a fixed width regardless of budget; costs round up, so a solution can never exceed B .

Two properties follow from the formulation rather than from added heuristics. **Demotion beats eviction:** a cheaper level with non-zero value dominates the skip option, so under pressure the optimiser first moves nodes down the ladder and only then drops them. And **the budget is a ceiling, not a target:** when the plan needs less, the surplus is simply unspent — a behaviour visible in §5.6, where the working set stops growing at 462 tokens no matter how much budget is offered.

3.6 The planner

Each turn: (i) classify the query; (ii) seed by hybrid lexical-vector search over *state nodes* — where the vector channel is a 256-dimensional hashing embedding over word and sub-word tokens rather than a learned semantic one, so the two channels are correlated rather than independent evidence, discarding seeds below a fraction of the top score; (iii) expand along dependency edges with depth and fan-out caps; (iv) assign each candidate its cheapest sufficient level from a routing table; (v) solve the knapsack; (vi) score the remainder for speculative prefetch.

Utility is a small, named, inspectable sum — similarity, graph proximity, membership of a demanded audit path, confidence, historical read-through, recency, a kind prior, a contradiction bonus, and penalties for staleness and for evidence whose dependents were superseded. We deliberately did not learn it. A wrong utility function is the design's dominant failure mode: it fills the window with cheap useless material *and looks efficient while doing so*. Keeping every term named means a bad answer can be traced to the term that caused it, and the implementation prints the per-node arithmetic on request.

3.7 Escalation

Deciding the cheapest level that can answer a query is the design's central open problem; there is no clean definition of sufficiency. We adopt the practical proxy: route by query type, then *measure* the error. If the compact form is insufficient the model replies `#ESCALATE <node_id>`, and the runtime re-plans with that node pinned at L0. Escalation rate is therefore a first-class metric — a rising rate is the signal that the router is wrong — and the extra tokens an escalation costs are charged to the turn that needed them.

3.8 Speculation and consistency

Nodes likely to be needed next are materialised in advance by an online logistic predictor over cheap features, trained on whether each prefetch was actually used. Prefetch consumes a *separate* materialisation budget denominated in build operations, never attention tokens: speculation that competes for the active window starves the computation it was meant to accelerate.

Finally, the Memory Runtime consolidates in the background while the Reasoner is thinking, and may invalidate a fact the Reasoner is mid-way through using. Every plan records a store version; after the model answers, the runtime checks whether anything in the working set was invalidated during the call and rebuilds rather than returning an answer grounded in state that no longer holds. This is snapshot isolation with an interrupt [12].

3.9 Authenticated durability

Every mechanism above assumes the memory runtime returns what it stored. That assumption does a great deal of work and nothing in the design so far earns it. A single JSON store can be edited in any text editor and reloaded without anything noticing; a disk can flip a bit in a span nobody has read for a year; a process killed mid-write can leave half an object. None of these are exotic, and all of them end the same way — the reasoner is handed something that is not what the runtime meant to store. A system whose central claim is that compressed state remains *auditable* has to be able to prove the state it audits is the state it wrote.

Context is therefore stored as a content-addressed, append-only container. Each object — document, span, node, edge — is addressed by the SHA-256 of its canonical encoding, so modifying an object changes its address and there is no such thing as an edited object, only a different one. The object set is summarised by a Merkle root, which lets any single object be verified against the committed state with $\log n$ hashes rather than by reading the store. Each generation is sealed by a checkpoint that chains to its parent, $\text{chain}_n = H(\text{chain}_{n-1} \parallel \text{checkpoint}_n)$, so altering any historical generation invalidates every later one. Generations are monotonic and the highest ever

accepted is persisted separately, which is secure-boot logic applied to memory state: an attacker cannot substitute an older, internally consistent version.

Two separations matter more than the mechanism. The first is that **integrity is not trust**. Cryptographic verification answers *was this modified since it was written*; it says nothing about whether what was written was true, or whether its author should have been believed. Those are a per-object trust label and an admission policy, kept deliberately distinct from verification state, and the gateway checks them *last* — so an object is shown to be intact, authentic and current before its content is judged, and a perfectly signed instruction from a low-trust source is still refused. The second is that detection is not repair. Signatures notice corruption and fix nothing, so a scrubber walks the store, recomputes digests, and repairs from a replica only once that replica verifies independently — hashing to the address being repaired *and* proving against the committed root. A scrubber that skips that check is a mechanism for spreading corruption evenly.

The whole admission path collapses to one rule, enforced in one place rather than checked in five and forgotten in a sixth: **no context enters trusted reasoning unless integrity is verified, the signature and provenance are valid, the generation is acceptable, the schema is valid, and policy allows it**.

What this does *not* provide is worth stating with equal force, because "cryptographically authenticated memory" invites a stronger reading than the construction supports. Without a signing key the container is **tamper-evident, not tamper-proof**. It detects bit rot, truncation, partial writes and casual edits, because none of those arrive with a recomputed chain. An attacker with write access to the directory can rewrite the objects, the chain, the manifest and the high-water mark together and produce a store that verifies cleanly, and no amount of hashing prevents that. Signatures raise the bar; genuine anti-rollback additionally requires the high-water mark to live somewhere the attacker cannot write. The implementation ships no signer at all (§4), so what is measured in §5.13 is evidence, not proof.

4 Implementation

The reference implementation is 13,721 lines of Rust in `src/` (edition 2024), plus roughly 3,000 lines of tests, with **no external crates**. Hashing, text scanning, JSON persistence, and argument parsing are all in-tree. The durability layer of §3.9 carries that rule to its honest limit: SHA-256 is implemented in-tree and checked against the published FIPS 180-4 vectors, because a hash is the one cryptographic primitive that can be verified against known output. Signatures and authenticated encryption are traits with *no bundled implementation* — hand-rolling Ed25519 or an AEAD, where the failure modes are constant-time

behaviour and nonce discipline rather than a wrong digest, would trade a dependency for a liability. A store therefore records honestly that it is unsigned rather than implying protection it does not have. This is a deliberate choice for a specification's reference implementation: a reader can audit every line that stands between the design and the measurement, and the build has no supply chain. A consequence worth stating is that the extractor is a hand-written scanner rather than a regular-expression engine — which turned out to read better as a parser than the pattern it replaced.

Table 1. Module map. Each design element has one home.

DESIGN ELEMENT	MODULE
L0 immutable raw store	<code>spans.rs</code>
Representation ladder	<code>ladder.rs</code>
Typed nodes, memory graph, provenance	<code>nodes.rs</code> , <code>graph.rs</code>
Node extraction, contradiction, supersession	<code>indexer.rs</code>
Hybrid lexical + vector index	<code>index.rs</code>
Level routing, escalation, de-escalation	<code>policy.rs</code>
Attention-budget knapsack	<code>budget.rs</code>
Relevance planner	<code>planner.rs</code>
Speculative prefetch	<code>speculation.rs</code>
Execution layer, memoisation	<code>execute.rs</code>
Telemetry	<code>telemetry.rs</code>
Runtime facade, persistence	<code>runtime.rs</code>

Determinism. Every ranking breaks ties on the item's own index and every persisted object preserves insertion order, so a plan never depends on hash iteration order. This matters more than it sounds: without it, an ablation cannot be attributed, because two runs of the same configuration may differ.

Structure. The graph is a vector of nodes plus an identifier index, with edges as index pairs — no reference-counted ownership cycles. Interior mutability is confined to exactly one place: the per-node level cache and the read counters, because materialising a cheaper representation is logically a read, and requiring an exclusive borrow there would prevent the planner from holding the graph while pricing what is in it.

Cost. The knapsack is $O(n \cdot B/q)$ for n candidates, budget B , and quantum q chosen to keep B/q at a fixed 400. Lexical retrieval is sub-linear through an inverted index. Vector retrieval has two paths: an exact scan over state nodes, and an approximate one using random-hyperplane LSH over the same hashing embedding — 8 tables of 12 bits, planes drawn from a fixed seed so that two runs of one configuration agree, multi-probe over 1-bit neighbours, and

a fall back to the exact scan whenever a query's buckets cannot supply k , so pruning is never a source of silent misses. The exact path is the default and every attention figure here comes from it; §6 reports what the approximate one costs and, more usefully, what it revealed.

The implementation ships 152 tests. One hundred and twenty-five cover the runtime described here, encoding its invariants as executable claims — that an unsourced fact is rejected, that a superseded node stays in the graph, that a stale node never enters a plan, that the stale-fact read metric can be driven non-zero and the guard drives it back to zero (§5.9), that disabling supersession makes the mutation probe's negative control fail rather than silently pass, that a hypothesis never renders like an observation, that the knapsack never exceeds its budget and matches a brute-force optimum on small instances, and that the working set stays bounded as history grows. The remaining twenty-seven cover a durability layer this report does not otherwise describe: a content-addressed object container with a hash chain over generations, in which editing history invalidates every generation after it, and bit-rot repair that is refused from any replica which does not itself verify.

5 Evaluation

5.1 The instrument

Evaluating a context runtime with a language model conflates two things: whether the right material was placed

Table 2. Probe set. Each targets a different way a long-context system fails.

PROBE	QUESTION	TESTS
corrected fact (mid-history)	what is the server ip?	supersession of a value corrected mid-transcript
corrected fact (late)	what is the deploy window?	supersession near the end
old fact, never repeated	who is the owner of the checkout service?	recall of an early fact outside any recent window
exact quote	quote the exact error message	byte-exact retrieval, not gist
justification	why did we roll back?	retrieval of a decision with its reason
detail buried in a long span	how many retry attempts were made?	a detail lost by compression; requires escalation
corrected fact (very late)	what is the hourly rate?	a correction 90% of the way through

5.3 Baselines

Two baselines share the instrument. **Full context** receives the entire transcript on every query. **Sliding window** receives the most recent 8,000 tokens. Both are given

Table 3. 300 turns, 27,362 tokens of history, B = 1,200, window = 8,000. Tokens are per query.

PROBE	FULL CONTEXT	SLIDING WINDOW	DCR	DCR TOKENS
corrected fact (mid-history)	ok	miss	ok	196
corrected fact (late)	ok	ok	ok	419
old fact, never repeated	miss	miss	ok	399
exact quote	miss	miss	ok	963

in the window, and whether the model then used it well. We therefore hold the reasoner fixed and *non-neural* across all conditions. Our instrument is a deterministic line matcher: it scores each line of whatever context it is given by token overlap with the question, returns the best line's value, and — when the best match is a compressed form it does not trust — emits an escalation request instead of guessing. It cannot paraphrase, cannot compose two facts, and cannot recover from being handed the wrong material.

This is a severe instrument, and that is the intent. If an answer is wrong, the runtime put the wrong thing in the window; no model quality is available to mask it. §6 states what the instrument correspondingly cannot show.

5.2 Corpus and probes

The corpus is a synthetic operations transcript: ten substantive documents followed by long-form noise — standup notes, triage sweeps, metrics digests, handovers — with three corrections injected at 35%, 80%, and 90% of the way through. The distractors are generated rather than written: 287 of the 300 documents are drawn from eight templates with an integer substituted, so each template recurs roughly thirty-six times. At 300 turns it is 27,362 tokens; the generator scales to any length. Seven probes exercise distinct failure modes (Table 2). Correctness is substring containment of a known answer, so it is checkable without a judge.

every advantage that is honest: on a scoring tie they prefer the *latest* matching line, so they answer a correction correctly whenever the correction is inside their view.

5.4 Result 1 — attention cost and context rot

PROBE	FULL CONTEXT	SLIDING WINDOW	DCR	DCR TOKENS
justification	ok	miss	ok	192
detail buried in a long span	ok	miss	ok	962
corrected fact (very late)	ok	ok	ok	102
correct	5 / 7	2 / 7	7 / 7	—
mean tokens per query	27,362	7,968	462	59× less

The sliding window's five misses are the important ones, because they are *structural*: the server-IP correction and the service owner both fall outside the most recent 8,000 tokens, and no model quality recovers a fact that is not in the input. The full-context column should be read as a floor rather than a ceiling — a real language model reading the whole transcript would beat a line matcher on these probes. What no model quality changes is the token column.

DCR's per-turn telemetry for this run: escalation rate 0.14, stale-fact read rate 0.00, audit-path completeness 1.00, three budget-forced demotions, zero budget overflows. The stale-fact read rate carries a caveat a zero always carries —

it is only informative if the instrument could have returned something else — which §5.9 discharges with an explicit positive control rather than asking the reader to take the zero on trust. The two expensive probes (963 and 962 tokens) are the two that escalated to raw bytes; that cost appears in the table rather than being hidden.

5.5 Result 2 — does k stay flat?

The design targets a per-turn cost of $O(k + r)$ — active state plus bounded fresh retrieval — while total history N grows without bound. The falsifiable version of that claim is simply: hold the budget fixed, grow the history, and watch the working set.




Figure 2. Total history and mean working-set size against corpus size, on log-spaced turn counts. The two panels carry their own scales deliberately — plotting a 283,000-token history and a 420-token working set on one axis would require a second scale and mislead. Both series are given exactly in Table 4.

Table 4. Scaling, B = 800. Correctness holds at every size, and is identical on the exact and approximate retrieval paths. The *ann* columns are a second runtime over the same corpus with LSH pruning enabled. Latencies are single-run and machine-dependent; the ratios are the durable part.

TURNS	HISTORY (TOK)	STATE NODES	MEAN K	MAX K	CORRECT	INGEST	QUERY	ANN QUERY	ANN K
100	8,482	124	416.7	764	7/7	0.02 s	0.6 ms	0.5 ms	416.7
300	27,362	298	411.6	787	7/7	0.05 s	1.4 ms	1.1 ms	451.7
1,000	93,442	911	398.6	793	7/7	0.22 s	3.3 ms	2.6 ms	355.6
3,000	283,253	2,661	413.3	795	7/7	1.02 s	12.6 ms	5.8 ms	346.9

Across a 33× growth in history the mean working set moves by 4% and the maximum by 4%, with correctness

unchanged. That range is bounded by the generator rather than by the runtime, and the bound matters more than the

number: this corpus emits **21 distinct documents at any size** — ten fixed documents, three corrections and eight noise templates, which a reader confirms by running the generator and grouping its output rather than by reading its source: thirteen shapes occur once and eight repeat, identically at 300 turns and at 30,000 — so growing it grows the transcript's *length* at fixed lexical *variety*. At 3,000 turns each of its eight distractor templates already appears 373 times, and at 30,000 turns 3,748 times. Flat *k* over that range shows the planner does not degrade with length at constant difficulty. It is not evidence that the working set

holds as the retrieval problem itself gets harder, and reporting a larger token count from this generator would overstate what had been asked of the system.

Table 4b answers the harder question on a second generator (bench --diverse) that draws distractors from four vocabularies by mixed-radix decomposition, giving 18,432 distinct documents rather than eight templates. It is added alongside the original and never replaces it: every other figure in this report comes from the standard corpus, and changing that corpus underneath the tables would invalidate them silently.

Table 4b. Scaling on the lexically diverse corpus, B = 1,200, approximate retrieval. *distinct* counts documents that differ other than by an embedded integer. The *ingest* and *query* columns are single runs on a shared machine and carry the spread described in §5.12 — up to 2× on cold timings and more on warm. They are reported for order of magnitude, not for reproduction; the deterministic columns are the ones worth quoting without conditions.

TURNS	HISTORY (TOK)	DISTINCT	STATE NODES	MEAN K	CORRECT	INGEST	QUERY
3,000	150,648	3,000	2,032	196	7/7	0.4 s	2 ms
10,000	512,964	10,000	6,762	220	7/7	4 s	5 ms
30,000	1,564,065	18,445	18,473	197	7/7	15 s	11 ms
80,000	4,191,322	18,445	48,651	235	7/7	178 s	33 ms

Across a 28× growth in history on this corpus the working set moves from 196 to 235 tokens, with correctness held at 7/7 throughout — at **4.19 million tokens of history and 48,651 state nodes**. Two limits are worth stating rather than leaving to be found. The generator exhausts its vocabulary at 18,432 combinations, so beyond roughly 18,000 turns the distinct count plateaus and documents begin to repeat: the 80,000-turn row averages about four copies of each distinct document, which is far better than 21 and is not unbounded diversity. And the probe set is the same seven questions at every size, so this measures whether a fixed retrieval task stays cheap as the haystack grows, not whether the working set holds as the number of things worth knowing grows.

The cost columns carry a result that revises §5.5. The standard corpus at 30,000 turns takes 226 s to ingest and 169 ms per query; the diverse corpus at the same document count takes 15 s and 11 ms. What survives the timing spread is the *ratio*, not the milliseconds: an order-of-magnitude gap between two corpora of identical document count, measured back to back under the same conditions, is not something a 2× run-to-run variance produces. Ingest cost on the standard generator is dominated by comparing each arriving document against thousands of near-duplicates, so the super-quadratic growth we first attributed to the linking passes is largely a property of the corpus. An inverted index built to fix it was both slower at scale and behaviour-changing — 457.1 tokens per query against 461.9 — and was reverted rather than published. The *ann* columns record something the latency argument alone would hide: approximate retrieval does not merely

find the same nodes faster, it finds a partly different set. At 300 turns it assembles a working set 10% larger than the exact path, and at 3,000 turns one 16% smaller. Correctness is unchanged at every size, and the implementation asserts that equality rather than assuming it — a claim worth one sentence of provenance, because until recently it was false. The check was a `debug_assert!`, and every benchmark here runs `--release`, where debug assertions are compiled out: a control that could not fire, sitting inside the table whose central claim it existed to guard. It is now a live assertion, verified by deliberately diverging the approximate path and confirming the run aborts. Even so, "same answers, fewer vectors scored" is a stronger claim than these four rows establish, and the divergence in *k* is where a larger probe set would look for the difference. This supports the bounded-attention claim. It does *not* establish $O(k + r)$ end to end: query latency still grows over the same range. We previously attributed that to vector retrieval being a linear scan over state nodes, and the attribution was wrong. Pruning with LSH removes roughly 96% of the vectors a query scores and takes the index call to a small fraction of a millisecond, and end-to-end latency at 3,000 turns falls by about a factor of three — but it does not flatten, because with the index that cheap the growth is in planning rather than retrieval. Sub-linear retrieval was a necessary condition and not the binding one; the gap between the design and the artefact is real and now sits somewhere else.

5.6 Result 3 — how small can the window be?

Table 5. Budget sweep at 300 turns. The planner takes what it needs, not what it is given.

B_ATTENTION	CORRECT	MEAN K	MAX K	ESCALATION RATE	DEMOTIONS
120	6 / 7	104.9	120	0.14	10
200	7 / 7	182.7	200	0.14	4
300	7 / 7	238.3	300	0.14	11
500	7 / 7	328.3	500	0.14	10
800	7 / 7	411.6	787	0.14	7
1,200	7 / 7	461.9	963	0.14	3
2,000	7 / 7	461.9	963	0.14	3
4,000	7 / 7	461.9	963	0.14	3

All seven probes are answerable from a **200-token** window over a 27,362-token history — a 137× reduction — and the curve is flat from there. Above 1,200 the numbers stop changing entirely: the plan is complete, and additional budget is left unspent. The failure at 120 tokens is graceful

and legible in the telemetry: demotions rise to 10 as the optimiser pushes nodes down the ladder before dropping any.

5.7 Result 4 — which mechanisms carry which probes?

Table 6. Ablation at B = 1,200, 300 turns. One mechanism disabled per row.

VARIANT	CORRECT	MEAN K	PROBES THAT FAIL
full runtime	7 / 7	461.9	—
no supersession	5 / 7	640.9	corrected fact (mid-history); corrected fact (late)
no escalation	6 / 7	447.1	detail buried in a long span
L2 only (no ladder)	5 / 7	381.4	exact quote; detail buried in a long span
no seed floor	7 / 7	638.9	— (costs 39% more tokens)
no reference linking	7 / 7	461.9	— (no effect on this corpus)
no graph expansion	7 / 7	242.0	— (and 46% <i>cheaper</i>)

Three mechanisms are load-bearing. Removing supersession costs two probes *and* 35% more tokens, because both sides of every correction stay live and compete for the window — the clearest demonstration we have that correction handling is what attacks context rot rather than merely tidying the store. Removing the ladder costs the two probes that need exact bytes. Removing escalation costs the buried-detail probe for a saving of 15 tokens per query — with a caveat about what that row measures, added after the fact and verified directly.

The *no escalation* row does not measure a reasoner that cannot ask. Setting the escalation budget to zero does not stop the reasoner emitting the request; it stops the runtime servicing it, and the unserved protocol token is then returned as the answer and graded against the expected substring. On the buried-detail probe the recorded answer is the literal string `#ESCALATE clai_e36d7d2dd2bb`, which contains no prose and therefore fails. So the row conflates "this probe cannot be answered without escalation" with "an escalation request was raised and nobody serviced it", and those are different claims. The session working the poor-fits list reports that the configuration the mechanism-loss claim actually describes — a harness with no way to emit the signal at all — scores

7/7; we have verified the token-grading directly and report their figure as theirs. Read the row as an upper bound on what escalation is worth, not a measurement of it.

Two results are negative, and we report them as such. **Reference linking** — deriving dependency edges from value mentions — changes nothing on this corpus, because our justification probe is answerable from a single node whose own text contains the reason. **Graph expansion** is worse than useless here: disabling it keeps all seven probes and saves 215 tokens per query, 46% of the working set. Neither result shows the mechanisms are wrong; it shows our corpus does not contain a question that requires composing two nodes, which is precisely the question those mechanisms exist to answer. A benchmark that cannot distinguish a graph from a flat store has not tested the graph. Building a probe set that isolates multi-hop retrieval is the first item of future work, and until it exists the honest position is that expansion is a cost we have not yet justified.

5.8 The price of the memory runtime

Ingestion of the 300-turn corpus produces 301 spans, 298 state nodes, and 705 edges, and builds 35 summaries and 823 vectors, in 0.04 s. That cost is paid once and amortised

over every subsequent query; it is not part of the per-turn token budget. Storage remains $O(N)$ and unbounded by design — the claim is bounded attention, not bounded storage.

5.9 A positive control for the stale-fact metric

A reviewer objected, correctly, that the `stale_fact_read_rate = 0.0` reported above is not yet a measurement.² A zero is only evidence if the run could have produced something else; otherwise it records that the exclusion path was never exercised, not that it works. A control that cannot fire and a control that passes are the same row in a table. We therefore add a control whose expected value is written down before it is run.

The setup poisons a derivation. A value is computed from an input fact; the input is then corrected, so the derivation depends on a superseded fact and is marked stale; a probe asks for exactly that value. We run it twice — once in production configuration, once with the staleness guard bypassed — with these expectations fixed in advance:

Table 7. Stale-fact positive control. Expected values fixed before execution.

RUN	EXPECTED RATE	MEASURED	VERDICT
guard on (production)	0.0	0.0	pass
guard off (bypassed)	1.0	1.0	pass

The bypassed run reaches 1.0, so the metric can return nonzero; the production run is driven to 0.0 by the guard.

Table 8. Read coverage against corpus size, $B = 800$, after replaying the fixed probe set.

TURNS	SPANS N	EVER SHOWN (L0)	COVERAGE	NEVER SEEN
100	101	10	9.9%	91
300	301	8	2.7%	293
1,000	1,001	6	0.6%	995
3,000	3,001	6	0.2%	2,995

The count of spans ever shown is flat — six to ten — while history grows thirty-fold, so coverage collapses from 9.9% to 0.2% and the unread region grows linearly with N . This is the dual cost of bounded attention, and it is exactly what the probe table in §5.4 cannot express: the runtime answers from compact facts without ever surfacing most of the history's content, so a span whose specifics silently governed an answer can never be caught by a probe. Storage grows $O(N)$; read coverage does not — the blind region does. Whether that region can be audited without replaying it — for instance by flagging never-covered spans whose facts feed admitted decisions — is left open (§7).

Table 9. Against baselines that retrieve rather than truncate. Same corpus, same reasoner, $B = 1,200$. Recursive is charged for every chunk it reads.

PROBE	FULL	RAG	SUMMARISE-ALL	RECURSIVE	DCR
-------	------	-----	---------------	-----------	-----

Only this pairing licenses reading the production zero in §5.4's telemetry as a guard that fired rather than one that was never called. Absent the control, that cell should not have been in the table.

5.10 Read coverage: the unconditioned dual

Every probe in §5.2 asks for something already known to be in the history. That is a conditioned sample, and the failure mode that motivates this work lies outside it: an answer comes back fluent precisely because nothing in the query revealed that a dropped detail governed it, so no one writes a probe for that detail. A benchmark can only sample the region a probe author already indexed; the silent failures live strictly outside it. No query-layer "did my context degrade" check escapes this, because any such check is conditioned on knowing what went missing.

One number is not so conditioned. `audit_path_completeness` is a property of answers; invert it and measure *read coverage* — the fraction of source spans that have ever appeared, at L0, in any assembled context. This replays no queries and is conditioned on nothing. We count L0 alone deliberately: it is the only ladder level that renders a span's actual bytes. Counting any admission overcounts badly, because corroboration collapses many agreeing spans behind one node — on this corpus a single L1 fact can sit on several hundred spans, none of whose specific content was shown. A span is covered only when its bytes reached the model.

5.11 Result 5 — a corpus built to make DCR lose

Every result so far comes from one corpus, and that corpus is kind to us. Its probes ask for material that similarity can find, which is why a plain top-k retriever does as well on it as full context. Table 9 makes that concrete: RAG over the same hybrid index answers 5/7, exactly matching the full-context baseline, at roughly 2.5× DCR's token cost. DCR's 7/7 comes from two probes specifically — a fact stated once and never repeated, where there is nothing for similarity to latch onto, and an exact quote, where top-k returns the right chunk but not the right bytes. That is a narrower claim than "7/7 versus 5/7" suggests, and stating it narrowly is the point.

PROBE	FULL	RAG	SUMMARISE-ALL	RECURSIVE	DCR
corrected fact (mid-history)	ok	ok	miss	ok	ok
corrected fact (late)	ok	ok	miss	miss	ok
old fact, never repeated	miss	miss	miss	miss	ok
exact quote	miss	miss	miss	miss	ok
justification	ok	ok	miss	ok	ok
detail buried in a long span	ok	ok	miss	ok	ok
corrected fact (very late)	ok	ok	ok	ok	ok
correct	5 / 7	5 / 7	1 / 7	4 / 7	7 / 7
mean tokens per query	27,362	1,168	1,197	31,074	462

Two of those columns are informative beyond the score. Summarising everything and truncating to the same budget collapses to 1/7: uniform compression destroys precisely the details the probes ask for, which is the strongest evidence here that *selective* representation is doing work that compression alone cannot. Recursive map-reduce costs more than full context — 31,074 tokens against 27,362 — because it reads every chunk and then reads its own digest. Unlimited coverage is not free, and charging honestly for it is what makes the comparison meaningful.

A benchmark that only rewards recall also rewards answering confidently at all times, which is the failure this architecture exists to avoid. Table 10 uses a second corpus built on the opposite principle: similarity is actively misleading, and two of the five probes are passed by *declining*. Multi-hop probes are scored on the assembled context rather than the answer, uniformly across every column, because the deterministic reasoner is a line matcher and cannot perform a join; scoring its output there would measure the stand-in model rather than the assembly.

Table 10. The discriminating corpus. Refusal probes are passed by *not* answering. DCR does not win this table.

PROBE	FULL	RAG	SUMMARISE-ALL	RECURSIVE	DCR
lexical decoy on a stale fact	ok	ok	miss	ok	ok
three-hop dependency (context)	ok	ok	miss	ok	ok
stale derivation (must refuse)	miss	miss	ok	ok	miss
absent fact (must refuse)	miss	miss	ok	miss	miss
contested fact (order decides)	miss	miss	miss	miss	miss
correct	2 / 5	2 / 5	2 / 5	3 / 5	2 / 5
mean tokens per query	27,382	1,133	1,180	31,180	237

DCR scores 2/5 and loses to recursive map-reduce. The three failures are the most useful output the evaluation has produced, and each names a real limit rather than a tuning problem.

Stale derivation. A figure that arrives as prose — “the incident cost estimate is 2,160 USD, computed from the engineer count, the incident hours and the hourly rate” — is served unchanged after the hourly rate is corrected. Invalidation propagates along dependency edges, and a number stated in text has none: the runtime only tracks derivations it computed itself. The fix is not obvious, because parsing “computed from X, Y and Z” into real edges would extend extraction into inference, which is exactly where wrong facts get cached with high confidence.

Absent fact. Asked for a value the history never states, DCR serves the adjacent line, where full context and RAG both

decline because their raw text scores below the reasoner's threshold. This is the cost of assembling eagerly: a window pre-filled with *selected* facts makes whatever is nearest look good. It is the mirror image of the failure §5.10 addresses — fluent because nothing in the query revealed that the answer was missing.

Contested fact. Two claims on one key disagree, neither is phrased as a correction, and the later silently supersedes the earlier with no marker reaching the window. This is documented behaviour, not a defect: ingest order is event order, and treating every plain restatement as a live contradiction would flood the window with disputes that are not disputes. But the cost is real and this row prices it — merge two transcripts and merge order decides what is true, silently.

One probe is worth reporting as a success for a different reason. The lexical decoy puts the query's exact words four times in the superseded document and once in the correction, so a system riding on retrieval attraction rather than the supersession edge answers with the stale value. DCR serves the correction, which is direct evidence that supersession is load-bearing rather than incidental.

Designing this corpus produced a methodological finding worth passing on. Three of the five probes initially failed for every condition, and for an uninteresting reason: their vocabulary collided with the filler text, so the probes measured the collision rather than the property they named. A probe that fails because the benchmark is broken is indistinguishable, in the table, from a probe that fails because the system is. Only inspecting the assembled window per probe separated them.

5.12 Result 6 – what rebuilding the workspace costs

Section 3.1 claims the working set can be destroyed and rebuilt from durable memory at any time. That is only a guarantee if rebuilding is cheap; if reconstruction requires re-deriving state, the invariant is nominal — true in principle, never exercised, and quietly false in the only situation where it matters. Rebuilding is one planning pass, not a replay: drop every cached representation, then classify, seed, expand, level-assign, solve, and materialise only what was admitted. The cost is therefore bounded by the working set rather than by history.

Measured over the seven probes at $B = 1,200$, a cold rebuild — every cached representation dropped, the window reassembled from raw spans alone — costs **1.55 ms on average against 0.15 ms warm** on an idle machine. The gap tracks how much L1 has to be rebuilt: probes that admit only cached facts rebuild nothing and cost the same either way, while the probe that pulls long spans rebuilds 21 summaries and dominates the mean, dominating the mean. At three orders of magnitude below a single model call, discarding the workspace between turns is a real option rather than a theoretical one.

These milliseconds are softer than they look, and we would rather say so than have a reader treat a slower reproduction as a regression. Re-running the same binary under load — two other processes compiling the same crate, load average 14 on 12 cores — moved the cold mean across 1.98, 2.73, 2.82, 3.05 and 4.18 ms and the warm mean from 0.15 to 1.02 ms, with the cold/warm ratio itself ranging from 4.6× to 12×. What survives that spread is the ordering and the shape: cold exceeds warm by roughly an order of magnitude, and the cost tracks how much L1 a probe forces to be rebuilt rather than how much history exists. The absolute figures are a single quiet-machine run and should be read as such.

The number does not license a stronger claim. Rebuild has two terms, and only one is bounded: materialisation is

capped by the attention budget, while retrieval is not (§5.5). The measurement should be read as "bounded by the working set, plus a retrieval term that is only now becoming sub-linear."

5.13 Result 7 – can the container detect tampering?

A security property that cannot fail a test is not a property, it is a claim. The durability layer of §3.9 is therefore exercised adversarially: corrupt a stored object, rewrite a historical checkpoint, and attempt a rollback to an older internally-consistent state. All three are detected — the first by the content address, the second by the chain, the third by the generation high-water mark — and a corrupted object additionally causes the runtime to refuse to load rather than to load partially, on the grounds that a memory with invisible gaps is worse than one that will not open.

The probe earned its place immediately. Its first run passed the rewritten-checkpoint case, because the chain digest covered only the Merkle root and the delta, leaving the timestamp, the object count and the policy hash editable without breaking anything. The chain now covers the whole checkpoint body. A later, exhaustive version of the same idea — mutate *every* field in turn and require each to break the chain — found a further gap that the single-case regression test could not have: the schema field was reconstructed from a constant when the digest was recomputed, so whatever was on disk never entered it.

What the result does not show is the case §3.9 already conceded: an attacker who rewrites objects, chain, manifest and high-water mark together produces a store that verifies. These three cases measure evidence of tampering, and nothing here measures resistance to it.

6 Threats to validity

We state these plainly because the results above are easy to over-read.

The corpus is synthetic. It was written to contain the failure modes the design targets, by the same people who designed it. It exercises correction, staleness, exact quoting, and detail buried by compression; it is not a sample of real agent traces, and its noise is not merely stylistically uniform but structurally repetitive: 287 distractor documents drawn from eight templates (§5.2) carry far less lexical variety than a real transcript. Part of the compression ratio in Table 3 therefore reflects that redundancy rather than the planner. The harder case — a corpus whose noise is topically close to its signal, which is the condition that most reliably induces context rot — is not covered here.

The instrument is not a language model. A line matcher cannot paraphrase or compose. This buys attribution — a wrong answer is always the runtime's fault — at the cost of external validity. In particular the baselines' correctness columns are *floors*: a competent model reading the full

transcript would answer more of these probes than our matcher does. Reading Table 3 as "DCR is more accurate than long context" would be a misreading. The defensible claims are the token counts and the sliding window's structural misses.

No RLM comparison. We motivate the design as a layer above recursive slicing of context, but we do not implement or measure against a recursive-language-model system. Any claim that DCR is faster or cheaper than RLM is therefore unsupported by this paper.

Tamper-evident is not tamper-proof. The durability layer detects corruption, historical rewriting, and rollback, and §5.13 measures that it does. It does not resist an adversary with write access to the store, who can rewrite objects, chain, manifest, and high-water mark together and produce a container that verifies. No signer ships with the implementation, so every result in §5.13 is evidence of tampering rather than resistance to it, and the anti-rollback mark is only as trustworthy as the medium holding it.

Controls are verified by argument, not by tooling. Several results here rest on controls: a positive control for the stale-fact metric (§5.9), a guard-fired column separating a respected supersession edge from a node that never surfaced, and the correctness-equality assertion behind Table 4. Four times a check has turned out not to be exercisable. Three were controls that could not fire: a metric that could not be exercised in the configuration that disabled it, a pass reachable through an unobserved shortcut, and an assertion compiled out of the release build that ran it. The fourth was the verification of the third — it fired correctly, on purpose, and kept no record a reader could repeat, so this paper briefly claimed a check that had to be taken on trust. That the fourth was committed while fixing the third is the useful part: describing this class confers no immunity to it. Five times an author held the relevant description in mind and produced a fresh instance anyway. Counting instances where the description was demonstrably available at the moment of the error, whether or not the error reached print: validating a control by confirming it passed, immediately after repairing one that could not fail; committing a verification no reader could repeat, while fixing an assertion that could not fire; asserting something a list read the same session contradicted; matching byte sequences instead of asking what a generator emits, while writing the note describing that exact substitution; and comparing a result against an unobserved counterfactual, while writing about claims that outrun their evidence. **The taxonomy is a diagnostic instrument, not a prophylactic one.** It identifies what went wrong once a second instrument has disagreed; it does not appear to prevent anything, and the guard is the instrument rather than the knowledge. We stop short of the stronger claim that naming a failure mode is worse than leaving it unnamed, because that needs a counterfactual we cannot

observe — only that, on this evidence, naming bought nothing. Each was found by someone asking a differently-shaped question about something already checked, and none by anyone re-running an existing check more carefully; we do not know how many remain. It is tempting to conclude that cross-review rather than tooling is what finds these, and a fifth instance is the counter-example: the editing script used to write this very entry called a string replacement without checking that its pattern matched, producing a successful-looking run against an unchanged file, and it was caught on the retry — by tooling, because the second attempt failed loudly where the first had failed silently. The narrower claim is the useful one: **tooling finds these exactly when it is built to fail noisily, and most tooling is not.** Two of the five would have been caught by that discipline alone. The advice this yields is also cheaper than seeking another reviewer, which is frequently unavailable: make the checks fail loudly. A narrow lint — no debug assertions in paths that only run under `--release` — would have caught the most recent at the moment it was written, and is the kind of check this evaluation still lacks. The general lesson is cheap: verifying a control means making it fail on purpose, because confirming that it passes is exactly what each broken version already did.

It is worth being precise about how far that generalises, because the obvious reading — that a second reviewer is indispensable and no amount of care substitutes for one — is not quite what the record shows. Four of the five were found by someone other than the author of the defect; the fifth was found by its own author, and not by re-reading it. It surfaced because a retry of the same edit failed loudly where the first attempt had failed silently, and every self-caught defect in this work has that shape: a re-measurement that disagreed with a stated tolerance, a count that disagreed with a prose number, two artefacts compared against each other that named different revisions. **Re-reading caught nothing, for anyone, at any point.** Recalling counted for less still: the last of these defects was a claim contradicted by a list in this very section, written from memory by an author who had read that list the same session. Memory is not a second instrument; it is a lossy copy of the first reading, and it is the copy that agrees with you. The operative requirement is a second *instrument*, not a second pair of eyes. Another person is one such instrument and a dependable one, but re-executing, re-measuring and counting are others, and they cost a lone author nothing. Stated as "you need a reviewer", the lesson is expensive and frequently unactionable; stated as "you need a reading that does not come from the same instrument as the first", it is neither.

One further observation bounds how much comfort to take from the test suite. The five were found, respectively, by a reader arguing on a public comment thread; by a probe an outside proposal specified; by an author deliberately sabotaging a code path to see whether a guard aborted; by

a retry of a silently-failed edit; and by a recount someone requested. **Not one was found by the test suite**, which was green before, during and after each of them, and is green now — as it was while stored objects were being re-addressed on every commit, and while the checkpoint chain covered only part of a checkpoint, both of which were found by running the command-line tool instead. This is not an argument against the tests: they encode the invariants and they catch regressions in them. It is the reason this section exists. A green suite is evidence about the properties somebody thought to encode and silent about the ones nobody did, which makes "152 tests pass" the weakest sentence available about this system. We would rather say so than let the count carry weight it has not earned.

A precision past the evidence is its own failure, and the worst-behaved one. The defects above divide into checks that could not fail and claims stated more generally than their evidence. A third kind appeared while this section was being written, and it is worth separating because it fails differently: the rebuild latencies in §5.12 were reported with an implied tolerance — "they move by 10–15% between runs" — that the quantity does not have. Re-running the same binary on a loaded machine moved the cold mean across a 2× range and the warm mean across nearly 7×. The measurement was sound and the sentence describing it was accurate; the invented part was the error bar around it. A claim past its evidence misleads a reader about *what* was shown. A precision past its evidence misleads about *how well*, and it is the only one of the three that turns a careful reader reproducing the work into someone who believes they have found a regression. Every timing in this report is now stated with its spread for that reason, and the deterministic columns — working-set size, hit counts, correctness — are the ones we would rather the argument rest on.

Three known failures, unfixed. Section 5.11 reports DCR losing: a derived figure that arrived as prose is served after its inputs are corrected, a question the history never answers gets an adjacent line rather than a refusal, and two disagreeing sources are settled silently by ingest order. These are reported rather than repaired because each has a principled fix that is larger than a patch — dependency extraction from prose, an admission floor, and a merge-order-independent conflict rule respectively — and shipping a narrow workaround for a benchmark row is how a benchmark stops measuring anything.

Two corpora, twelve probes. Twelve probes on two generators is enough to falsify a claim, not enough to establish a distribution — and both generators were written by the same people who designed the system. The ablation's two negative results, and the three failures in §5.11, are direct evidence of the probe sets' limits rather than of their adequacy.

Extraction is conservative and non-neural. The scanner proposes state only for unambiguous shapes. Prose-heavy material stays at L0 and is found by search instead. A model-driven extractor would find more state — and would introduce the failure mode this design most fears, a confidently cached wrong claim, which mandatory provenance mitigates but does not solve.

Token counts are estimated by a character heuristic rather than a real tokeniser, uniformly across all conditions. Ratios are therefore more trustworthy than absolute counts.

Retrieval is no longer the binding constraint. The cost model requires sub-linear retrieval as one of its three conditions; the approximate index supplies it, and end-to-end latency still grows, because the planner's own cost does and we have not shown that to be sub-linear. The bounded-attention result stands. The end-to-end complexity claim does not, and the reason it does not has changed — worth saying plainly, because an earlier version of this section named the wrong cause with more confidence than the evidence supported.

7 Open problems

Defining sufficiency. $\text{sufficient}(L, q)$ has no clean definition. We route by query type and measure escalation, which is a proxy with a measurable error rate, not a theory. A learned router with escalation rate as its loss is the obvious next step.

Estimating utility. Our $U(x)$ is hand-weighted. Learning it from downstream answer quality is attractive and dangerous in equal measure: the failure mode of a learned utility is a window full of cheap material that scores well and answers nothing, and the telemetry that would catch it is exactly the escalation rate the router is already using.

Trusting extracted state. Node extraction is the point where a wrong fact enters with high confidence. Mandatory spans, confidence thresholds, and contradiction edges reduce the blast radius. What error rate is tolerable, and how it should be measured, is unresolved.

Sub-linear retrieval over state. An approximate-nearest-neighbour index behind the vector interface would close the gap identified in §5.5. The interesting version of the question is whether graph structure can replace part of the vector search rather than merely accelerating it.

Multi-hop evaluation. The probe set needs questions whose answers exist only as a composition of two nodes, so that graph expansion and reference linking can be validated or discarded on evidence.

Auditing the unread region. Read coverage (§5.10) names where confident wrong answers can hide but does not detect them: a never-covered span is only dangerous if its content would have changed an answer, and deciding that without surfacing the span is unsolved. A cheap proxy —

flag never-covered spans whose extracted facts feed an admitted decision — is worth testing, but a span the extractor never turned into a fact is invisible even to that.

Bounded revalidation. Cascades are bounded and the tail is deferred, but we do not characterise the staleness that deferral admits, nor the cost of rebuilding a workspace from the store — the invariant that the working set is disposable is only useful if rebuilding is cheap.

8 Conclusion

Context rot is usually treated as a prompt-engineering problem, addressed by summarising old turns when someone remembers to. We have argued it is a systems problem with a known shape: unbounded external state, an explicit budget on what may be attended to, and an optimiser that decides which representation of which fact is worth its tokens this turn. Framed that way, most of the machinery already exists in database systems — provenance, view maintenance, snapshot isolation, cache admission — and can be imported rather than reinvented.

The measurements support a narrow, useful claim: a working set of a few hundred tokens can answer the same questions as a 27,000-token transcript, and stays the same size as that transcript grows 33× larger. They also show where the design is not yet earned — retrieval is still linear, and two of our mechanisms are not validated by our own benchmark. We would rather publish that than a table with no negative rows.

Availability

Every figure in this paper was measured at revision 8b7162f. Benchmark numbers move whenever retrieval, extraction or the planner moves, so they carry a revision rather than a date; re-running the commands below on a later revision will not reproduce them exactly, and the repository records what to re-measure. The specification wiki, the Rust implementation, and every experiment in this paper are available at github.com/s-b-repo/subnext. All results reproduce offline with no API key and no network access:

```
cargo run --release -- bench
# Table 3
cargo run --release -- bench --scaling
# Table 4, Figure 2
cargo run --release -- bench --sweep
# Table 5
cargo run --release -- bench --ablate
# Table 6
cargo run --release -- bench --poison
# Table 7 (stale-fact control)
cargo run --release -- bench --coverage
# Table 8 (read coverage, and pair coverage)
cargo run --release -- bench --mutate
# mutation and correction, plain and adversarial
cargo run --release -- bench --multihop
# does expansion buy anything on a join?
cargo run --release -- bench --decay
# recency prefilter: latency against recall
cargo run --release -- bench --consolidate
# a write landing mid-turn
cargo run --release -- bench --diverse
# Table 4b, to 4.19M tokens
cargo run --release -- bench --baselines
# Tables 9 and 10
cargo run --release -- bench --rebuild
# §5.12
cargo run --release -- bench --tamper
# §5.13
cargo test
# 152 tests
```

References

1. A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin. Attention is all you need. *Advances in Neural Information Processing Systems*, 2017.
2. N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: how language models use long contexts. *Transactions of the ACL*, 2024.
3. P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W. Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 2020.
4. I. Beltagy, M. E. Peters, and A. Cohan. Longformer: the long-document transformer. *arXiv:2004.05150*, 2020.
5. T. Dao, D. Y. Fu, S. Ermon, A. Rudra, and C. Ré. FlashAttention: fast and memory-efficient exact attention with IO-awareness. *Advances in Neural Information Processing Systems*, 2022.
6. S. Robertson and H. Zaragoza. The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4), 2009.
7. V. Karpukhin, B. Oğuz, S. Min, P. Lewis, L. Wu, S. Edunov, D. Chen, and W. Yih. Dense passage retrieval for open-domain question answering. *EMNLP*, 2020.
8. J. S. Park, J. C. O'Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein. Generative agents: interactive simulacra of human behavior. *UIST*, 2023.
9. C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez. MemGPT: towards LLMs as operating systems. *arXiv:2310.08560*, 2023.
10. P. Buneman, S. Khanna, and W. C. Tan. Why and where: a characterization of data provenance. *International Conference on Database Theory*, 2001.
11. A. Gupta and I. S. Mumick. Maintenance of materialized views: problems, techniques, and applications. *IEEE Data Engineering Bulletin*, 18(2), 1995.
12. H. Berenson, P. Bernstein, J. Gray, J. Melton, E. O'Neil, and P. O'Neil. A critique of ANSI SQL isolation levels. *ACM SIGMOD*, 1995.
13. H. Kellerer, U. Pferschy, and D. Pisinger. *Knapsack Problems*. Springer, 2004.
14. J. E. Smith. A study of branch prediction strategies. *International Symposium on Computer Architecture*, 1981.

¹ The recursive-language-model framing used here is taken from the Dynamic Context Runtime specification wiki (github.com/s-b-repo/subnext), which describes the primitive as treating context as a variable in a REPL that the model can slice and recursively query. This paper responds to that framing; it does not evaluate against any particular implementation of it.

² The positive control of §5.9 and the read-coverage metric of §5.10 were prompted by a public review from the commenter `vespermind`, who argued that the stale-fact zero was unmeasured and that read coverage is the one degradation signal not conditioned on knowing what went missing. Both points were correct; the implementation of read coverage also surfaced the corroboration over-count noted in §5.10, which the L0-only definition corrects.